

Note on Bandit Algorithms

September 4, 2026

1 Bandits, Probability and Concentration

1.1 Introduction

Definition.

- Actions : $A_1, A_2, \dots \in \mathcal{A}$ and we call $\mathcal{A}$ the action set.
- Rewards : $X_1, X_2, \dots \in \mathcal{R}$.
- History : $H_{t-1} = (A_1, X_1, \dots, A_{t-1}, X_{t-1})$.
- Policy : $\pi \in \Pi : H_{t-1} \mapsto A_t$ and we call Π the competitor class.
- Environment: $\mathcal{E} : A_t \mapsto X_t$.
- Regret of a learner relative to a policy π until $n \in \mathbb{N}$ is

$$R_n = \mathbb{E} \left[\sum_{t=1}^n \mathcal{E}(\pi(H_{t-1})) - \sum_{t=1}^n X_t \right]. \quad (1)$$

where H_0 represents the initial condition.

Example. An environment is called a **stochastic Bernoulli bandit** if

- $\mathcal{A} = \{1, 2, \dots, k\}$,
- $X_t \in \{0, 1\}$,
- $\forall a \in \mathcal{A}, \exists \mu_a \in [0, 1]^k$ such that $\mathbb{P}[X_t = 1 | A_t = a] = \mu_a$.

Here the optimal policy is to play fixed action $a^* = \operatorname{argmax}_{a \in \mathcal{A}} \mu_a$. That is, the competitor class Π is the set of all constant policies $\Pi = \{\pi_1, \dots, \pi_k\}$ where π_i chooses action i in every round. The regret relative to the optimal policy π^* until n rounds is

$$\begin{aligned} R_n &= \mathbb{E} \left[\sum_{t=1}^n \mathcal{E}(\pi^*(H_{t-1})) - \sum_{t=1}^n X_t \right] \\ &= n \max_{a \in \mathcal{A}} \mu_a - \mathbb{E} \left[\sum_{t=1}^n X_t \right]. \end{aligned}$$

Definition. Stochastic stationary bandits are the environments restricted to generate the rewards in response to each action from a distribution independent of the previous history.

Example (Linear Stochastic Bandits).

- $\mathcal{A} \subset \mathbb{R}^d$,
- $X_t = \langle a, \theta \rangle + \eta_t$ for $\theta \in \mathbb{R}^d, a \in \mathcal{A}$ and η_t is a standard Gaussian.

1.2 Foundations of Probability

Definition. Let an event set Ω be given.

- A set $\mathcal{F} \subset 2^\Omega$ is called a **σ -algebra** on Ω if
 - (1) $\Omega \in \mathcal{F}$,
 - (2) $A \in \mathcal{F} \implies A^c \in \mathcal{F}$,
 - (3) $A_1, A_2, \dots \in \mathcal{F} \implies \bigcup_{i=1}^\infty A_i \in \mathcal{F}$.
- $\mathcal{G} \subseteq \mathcal{F}$ is a **sub- σ -algebra** if $\mathcal{G}$ is a σ -algebra on Ω .
- $(\Omega, \mathcal{F})$ is called a **measurable space** if $\mathcal{F}$ is a σ -algebra on Ω .
- $\mathcal{B}(\mathbb{R})$ is the **Borel σ -algebra** on $\mathbb{R}$, which is the smallest σ -algebra containing all open intervals (a, b) for $a, b \in \mathbb{R}$.

Definition. Let a measurable space $(\Omega, \mathcal{F})$ be given.

- A function $\mu : \mathcal{F} \mapsto [0, \infty]$ is **σ -additive** if
 - (1) $\mu(\emptyset) = 0$,
 - (2) $A_1, A_2, \dots \in \mathcal{F}$ are disjoint $\implies \mu(\bigcup_{i=1}^\infty A_i) = \sum_{i=1}^\infty \mu(A_i)$.
- A function $\mu : \mathcal{F} \mapsto [0, \infty]$ is a **measure** if μ is σ -additive.
- $(\Omega, \mathcal{F}, \mu)$ is called a **measure space** if μ is a measure on $(\Omega, \mathcal{F})$.
- A measure μ is a **probability measure** if $\mu(\Omega) = 1$.
- $(\Omega, \mathcal{F}, \mu)$ is called a **probability space** if μ is a probability measure on $(\Omega, \mathcal{F})$.
- A function $X : \Omega \mapsto \mathbb{R}$ is called a **measurable function** if $\forall B \in \mathcal{B}(\mathbb{R}), X^{-1}(B) \in \mathcal{F}$.
- A **filtration** is a sequence $(\mathcal{F}_t)_{t=0}^n$ of sub- σ -algebras of $\mathcal{F}$ where $\mathcal{F}_t \subseteq \mathcal{F}_{t+1}$ for all $t < n$.

Definition. Let a measure space $(\Omega, \mathcal{F}, \mu)$ be given.

- A function $f : \Omega \mapsto [0, \infty]$ is called a **simple function** if $f = \sum_{i=1}^n a_i \mathbb{1}_{A_i}$ for some $a_1, a_2, \dots, a_n \in [0, \infty]$ and $A_1, A_2, \dots, A_n \in \mathcal{F}$.
- The **integral of a simple function** $f = \sum_{i=1}^n a_i \mathbb{1}_{A_i}$ is defined as $\int_\Omega f d\mu := \sum_{i=1}^n a_i \mu(A_i)$.
- The **integral of a non-negative measurable function** $f : \Omega \mapsto [0, \infty]$ is defined as $\int_\Omega f d\mu := \sup_g \int_\Omega g d\mu$ where the supremum is taken over all simple functions g such that $0 \leq g(\omega) \leq f(\omega)$ for all $\omega \in \Omega$.
- The **integral of a measurable function** $f : \Omega \mapsto \mathbb{R}$ is defined as $\int_\Omega f d\mu := \int_\Omega f^+ d\mu - \int_\Omega f^- d\mu$ where $f^+ = \max(f, 0)$ and $f^- = \max(-f, 0)$.
- A function $f : \Omega \mapsto \mathbb{R}$ is called **integrable** if $\int_\Omega |f| d\mu < \infty$.

Definition. Let a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ be given.

- A measurable function $X : \Omega \mapsto \mathbb{R}$ is called a **random variable**.
- The **conditional probability** $\mathbb{P}[A|B] := \frac{\mathbb{P}[A \cap B]}{\mathbb{P}[B]}$ for $A, B \in \mathcal{F}$ and $\mathbb{P}[B] > 0$.
- Two events $A, B \in \mathcal{F}$ are called **independent** if $\mathbb{P}[A \cap B] = \mathbb{P}[A] \cdot \mathbb{P}[B]$.
- Two random variables X, Y are called **independent** if $\forall A, B \in \mathcal{B}(\mathbb{R}), \mathbb{P}[X \in A, Y \in B] = \mathbb{P}[X \in A] \cdot \mathbb{P}[Y \in B]$.

Definition. Let a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ and a random variable X be given.

- The **expectation** of a random variable X is defined as $\mathbb{E}[X] := \int_{\Omega} X d\mathbb{P}$.
- For a sub- σ -algebra $\mathcal{H} \subseteq \mathcal{F}$, the **conditional expectation** of X given $\mathcal{H}$ is denoted by $\mathbb{E}[X|\mathcal{H}]$ and is defined as a random variable g on Ω such that for all $H \in \mathcal{H}$,

$$\int_H g d\mathbb{P} = \int_H X d\mathbb{P}.$$

Theorem (2.11). If $(\Omega, \mathcal{F}, \mathbb{P})$ is a probability space and X is a integrable random variable and $\mathcal{H} \subseteq \mathcal{F}$ is a sub- σ -algebra, there exists a unique (almost surely) random variable $\mathbb{E}[X|\mathcal{H}]$ such that for all $H \in \mathcal{H}$,

$$\int_H \mathbb{E}[X|\mathcal{H}] d\mathbb{P} = \int_H X d\mathbb{P}.$$

Theorem (2.12). Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space, $\mathcal{G}, \mathcal{G}_1, \mathcal{G}_2 \subseteq \mathcal{F}$ be sub- σ -algebras of $\mathcal{F}$, X, Y be integrable random variables on $(\Omega, \mathcal{F}, \mathbb{P})$. Then the following properties hold:

- (1) If $X \geq 0$ almost surely, then $\mathbb{E}[X|\mathcal{G}] \geq 0$ almost surely.
- (2) $\mathbb{E}[1|\mathcal{G}] = 1$ almost surely.
- (3) $\mathbb{E}[X + Y|\mathcal{G}] = \mathbb{E}[X|\mathcal{G}] + \mathbb{E}[Y|\mathcal{G}]$ almost surely.
- (4) If Y is $\mathcal{G}$ -measurable and $\mathbb{E}[XY] < \infty$, then $\mathbb{E}[XY|\mathcal{G}] = Y\mathbb{E}[X|\mathcal{G}]$ almost surely.
- (5) If $\mathcal{G}_1 \subseteq \mathcal{G}_2$, then $\mathbb{E}[X|\mathcal{G}_1] = \mathbb{E}[\mathbb{E}[X|\mathcal{G}_2]|\mathcal{G}_1]$ almost surely.

1.3 Stochastic Processes and Markov Chains

Definition. Let two measurable spaces $(\mathcal{X}, \mathcal{F})$ and $(\mathcal{Y}, \mathcal{G})$ be given.

- $(\mathcal{X}, \mathcal{F})$ and $(\mathcal{Y}, \mathcal{G})$ are called **isomorphic** if there exists a bijective function $f : \mathcal{X} \mapsto \mathcal{Y}$ such that $A \in \mathcal{F} \iff f(A) \in \mathcal{G}$.
- $(\mathcal{X}, \mathcal{F})$ is a **Borel space** if it is isomorphic to $(\mathbb{R}, \mathcal{B}(A))$ for some set $A \in \mathcal{B}(\mathbb{R})$.

Theorem (3.1). Let μ be a probability measure on a Borel space S and λ be the Lebesgue measure on $([0, 1], \mathcal{B}([0, 1]))$. Then there exists a sequence of independent random element $X_1, X_2, \dots$ on $([0, 1], \mathcal{B}([0, 1]), \lambda)$ such that $\lambda_{X_t}(\cdot) := \lambda(X_t^{-1}(\cdot)) = \mu(\cdot)$ for all t .

Definition. For "countable" set $\mathcal{T}$, A stochastic process is a collection of random variables $X_t : t \in \mathcal{T}$.

Theorem (3.2). Let Borel spaces $(\Omega_n, \mathcal{F}_n)$ and measures μ_n on $(\Omega_1 \times \dots \times \Omega_n, \mathcal{F}_1 \otimes \dots \otimes \mathcal{F}_n)$ are given for each $n \in \mathbb{N}$.

$$\begin{aligned} &\forall n \in \mathbb{N} \text{ and } A \in \mathcal{F}_1 \otimes \dots \otimes \mathcal{F}_n, \quad \mu_{n+1}(A \times \Omega_{n+1}) = \mu_n(A) \\ &\implies \exists \text{ a probability space } (\Omega, \mathcal{F}, \mathbb{P}) \\ &\text{and random elements } X_1, X_2, \dots \text{ with } X_t : \Omega \rightarrow \Omega_t \text{ such that } \mathbb{P}_{X_1, \dots, X_n} = \mu_n \text{ for all } n. \end{aligned}$$

Definition. Let $(\mathcal{X}, \mathcal{F})$ and $(\mathcal{Y}, \mathcal{G})$ be measurable spaces. A **probability kernel** or **Markov kernel** from $(\mathcal{X}, \mathcal{F})$ to $(\mathcal{Y}, \mathcal{G})$ is a function $K : \mathcal{X} \times \mathcal{G} \rightarrow [0, 1]$ such that

- (a) $K(x, \cdot)$ is a measure for all $x \in \mathcal{X}$ and
- (b) $K(\cdot, A)$ is $\mathcal{F}$ -measurable for all $A \in \mathcal{G}$.

Notation: $K_x(A)$ or $K(A|x)$

Definition. Let K_1 be a probability kernel from $(\mathcal{X}, \mathcal{F}) \rightarrow (\mathcal{Y}, \mathcal{G})$ and K_2 be a probability kernel from $(\mathcal{Y}, \mathcal{G}) \rightarrow (\mathcal{Z}, \mathcal{H})$. Then the product kernel $K_1 \otimes K_2$ is the probability kernel from $(\mathcal{X}, \mathcal{F})$ to $(\mathcal{Y} \times \mathcal{Z}, \mathcal{G} \otimes \mathcal{H})$ defined by

$$(K_1 \otimes K_2)(x, A) = \int_{\mathcal{Y}} \int_{\mathcal{Z}} \mathbb{1}_A((y, z)) K_2(y, dz) K_1(x, dy).$$

Theorem (3.3). Let $(\Omega_n, \mathcal{F}_n)_{n=1}^{\infty}$ be a sequence of measurable spaces.

K_1 : a probability measure on $(\Omega_1, \mathcal{F}_1)$,

K_n : a probability kernel from $\Pi_{t=1}^{n-1} \Omega_t$ to Ω_n for $n \geq 2$.

$\implies \exists$ a probability space $(\Omega, \mathcal{F}, \mathbb{P})$

and random elements $(X_t)_{t=1}^{\infty}$ with $X_t : \Omega \rightarrow \Omega_t$ such that $\mathbb{P}_{X_1, \dots, X_n} = \bigotimes_{t=1}^n K_t$ for all n .

Definition. A **homogeneous Markov chain** is a sequence of random elements $(X_t)_{t=1}^{\infty}$ taking values in state space $\mathcal{S} = (\mathcal{X}, \mathcal{F})$ and with

$$\mathbb{P}(X_{t+1} \in \cdot | X_1, \dots, X_t) = \mathbb{P}(X_{t+1} \in \cdot | X_t) = \mu(X_t, \cdot), \quad \text{a.s.}$$

Here μ is a probability kernel from $(\mathcal{X}, \mathcal{F})$ to $(\mathcal{X}, \mathcal{F})$.

Definition. A $\mathbb{F}$ -adapted sequence of random variables $(X_t)_{t \in \mathbb{N}_+}$ (i.e. X_t is $\mathcal{F}_t$ -measurable for all $t \in 1, \dots, n$) is a $\mathbb{F}$ -adapted **martingale** if

(a)

(b)

1.4 Stochastic Bandits

Definition. A **stochastic bandit** is a collection of distributions $\nu = \{P_a : a \in \mathcal{A}\}$.

Here we are challenging to maximise the total reward $S_n = \sum_{t=1}^n X_t$ over n rounds.