

Note on Bandit Algorithms

September 30, 2026

1 Bandits, Probability and Concentration

1.1 Introduction

Definition.

- Actions : $A_1, A_2, \dots \in \mathcal{A}$ and we call $\mathcal{A}$ the action set.
- Rewards : $X_1, X_2, \dots \in \mathcal{R}$.
- History : $H_{t-1} = (A_1, X_1, \dots, A_{t-1}, X_{t-1})$.
- Policy : $\pi \in \Pi : H_{t-1} \mapsto A_t$ and we call Π the competitor class.
- Environment: $\mathcal{E} : A_t \mapsto X_t$.
- Regret of a learner relative to a policy π until $n \in \mathbb{N}$ is

$$R_n = \mathbb{E} \left[\sum_{t=1}^n \mathcal{E}(\pi(H_{t-1})) - \sum_{t=1}^n X_t \right].$$

where H_0 represents the initial condition.

Example A. n environment is called a **stochastic Bernoulli bandit** if

- $\mathcal{A} = \{1, 2, \dots, k\}$,
- $X_t \in \{0, 1\}$,
- $\forall a \in \mathcal{A}, \exists \mu_a \in [0, 1]^k$ such that $\mathbb{P}[X_t = 1 | A_t = a] = \mu_a$.

Here the optimal policy is to play fixed action $a^* = \operatorname{argmax}_{a \in \mathcal{A}} \mu_a$. That is, the competitor class Π is the set of all constant policies $\Pi = \{\pi_1, \dots, \pi_k\}$ where π_i chooses action i in every round. The regret relative to the optimal policy π^* until n rounds is

$$\begin{aligned} R_n &= \mathbb{E} \left[\sum_{t=1}^n \mathcal{E}(\pi^*(H_{t-1})) - \sum_{t=1}^n X_t \right] \\ &= n \max_{a \in \mathcal{A}} \mu_a - \mathbb{E} \left[\sum_{t=1}^n X_t \right]. \end{aligned}$$

Definition. Stochastic stationary bandits are the environments restricted to generate the rewards in response to each action from a distribution independent of the previous history.

Example Linear Stochastic Bandits.

- $\mathcal{A} \subset \mathbb{R}^d$,
- $X_t = \langle a, \theta \rangle + \eta_t$ for $\theta \in \mathbb{R}^d, a \in \mathcal{A}$ and η_t is a standard Gaussian.

1.2 Foundations of Probability

Definition. Let an event set Ω be given.

- A set $\mathcal{F} \subset 2^\Omega$ is called a **σ -algebra** on Ω if
 - (1) $\Omega \in \mathcal{F}$,
 - (2) $A \in \mathcal{F} \implies A^c \in \mathcal{F}$,
 - (3) $A_1, A_2, \dots \in \mathcal{F} \implies \bigcup_{i=1}^\infty A_i \in \mathcal{F}$.
- $\mathcal{G} \subseteq \mathcal{F}$ is a **sub- σ -algebra** if $\mathcal{G}$ is a σ -algebra on Ω .
- $(\Omega, \mathcal{F})$ is called a **measurable space** if $\mathcal{F}$ is a σ -algebra on Ω .
- $\mathcal{B}(\mathbb{R})$ is the **Borel σ -algebra** on $\mathbb{R}$, which is the smallest σ -algebra containing all open intervals (a, b) for $a, b \in \mathbb{R}$.

Definition. Let a measurable space $(\Omega, \mathcal{F})$ be given.

- A function $\mu : \mathcal{F} \mapsto [0, \infty]$ is **σ -additive** if
 - (1) $\mu(\emptyset) = 0$,
 - (2) $A_1, A_2, \dots \in \mathcal{F}$ are disjoint $\implies \mu(\bigcup_{i=1}^\infty A_i) = \sum_{i=1}^\infty \mu(A_i)$.
- A function $\mu : \mathcal{F} \mapsto [0, \infty]$ is a **measure** if μ is σ -additive.
- $(\Omega, \mathcal{F}, \mu)$ is called a **measure space** if μ is a measure on $(\Omega, \mathcal{F})$.
- A measure μ is a **probability measure** if $\mu(\Omega) = 1$.
- $(\Omega, \mathcal{F}, \mu)$ is called a **probability space** if μ is a probability measure on $(\Omega, \mathcal{F})$.
- A function $X : \Omega \mapsto \mathbb{R}$ is called a **measurable function** if $\forall B \in \mathcal{B}(\mathbb{R}), X^{-1}(B) \in \mathcal{F}$.
- A **filtration** is a sequence $(\mathcal{F}_t)_{t=0}^n$ of sub- σ -algebras of $\mathcal{F}$ where $\mathcal{F}_t \subseteq \mathcal{F}_{t+1}$ for all $t < n$.

Definition. Let a measure space $(\Omega, \mathcal{F}, \mu)$ be given.

- A function $f : \Omega \mapsto [0, \infty]$ is called a **simple function** if $f = \sum_{i=1}^n a_i \mathbb{1}_{A_i}$ for some $a_1, a_2, \dots, a_n \in [0, \infty]$ and $A_1, A_2, \dots, A_n \in \mathcal{F}$.
- The **integral of a simple function** $f = \sum_{i=1}^n a_i \mathbb{1}_{A_i}$ is defined as $\int_\Omega f d\mu := \sum_{i=1}^n a_i \mu(A_i)$.
- The **integral of a non-negative measurable function** $f : \Omega \mapsto [0, \infty]$ is defined as $\int_\Omega f d\mu := \sup_g \int_\Omega g d\mu$ where the supremum is taken over all simple functions g such that $0 \leq g(\omega) \leq f(\omega)$ for all $\omega \in \Omega$.
- The **integral of a measurable function** $f : \Omega \mapsto \mathbb{R}$ is defined as $\int_\Omega f d\mu := \int_\Omega f^+ d\mu - \int_\Omega f^- d\mu$ where $f^+ = \max(f, 0)$ and $f^- = \max(-f, 0)$.
- A function $f : \Omega \mapsto \mathbb{R}$ is called **integrable** if $\int_\Omega |f| d\mu < \infty$.

Definition. Let a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ be given.

- A measurable function $X : \Omega \mapsto \mathbb{R}$ is called a **random variable**.
- The **conditional probability** $\mathbb{P}[A|B] := \frac{\mathbb{P}[A \cap B]}{\mathbb{P}[B]}$ for $A, B \in \mathcal{F}$ and $\mathbb{P}[B] > 0$.
- Two events $A, B \in \mathcal{F}$ are called **independent** if $\mathbb{P}[A \cap B] = \mathbb{P}[A] \cdot \mathbb{P}[B]$.
- Two random variables X, Y are called **independent** if $\forall A, B \in \mathcal{B}(\mathbb{R}), \mathbb{P}[X \in A, Y \in B] = \mathbb{P}[X \in A] \cdot \mathbb{P}[Y \in B]$.

Definition. Let a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ and a random variable X be given.

- The **expectation** of a random variable X is defined as $\mathbb{E}[X] := \int_{\Omega} X d\mathbb{P}$.
- For a sub- σ -algebra $\mathcal{H} \subseteq \mathcal{F}$, the **conditional expectation** of X given $\mathcal{H}$ is denoted by $\mathbb{E}[X|\mathcal{H}]$ and is defined as a random variable g on Ω such that for all $H \in \mathcal{H}$,

$$\int_H g d\mathbb{P} = \int_H X d\mathbb{P}.$$

Theorem 2.11. If $(\Omega, \mathcal{F}, \mathbb{P})$ is a probability space and X is an integrable random variable and $\mathcal{H} \subseteq \mathcal{F}$ is a sub- σ -algebra, there exists a unique (almost surely) random variable $\mathbb{E}[X|\mathcal{H}]$ such that for all $H \in \mathcal{H}$,

$$\int_H \mathbb{E}[X|\mathcal{H}] d\mathbb{P} = \int_H X d\mathbb{P}.$$

Theorem 2.12. Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space, $\mathcal{G}, \mathcal{G}_1, \mathcal{G}_2 \subseteq \mathcal{F}$ be sub- σ -algebras of $\mathcal{F}$, X, Y be integrable random variables on $(\Omega, \mathcal{F}, \mathbb{P})$. Then the following properties hold:

- (1) If $X \geq 0$ almost surely, then $\mathbb{E}[X|\mathcal{G}] \geq 0$ almost surely.
- (2) $\mathbb{E}[1|\mathcal{G}] = 1$ almost surely.
- (3) $\mathbb{E}[X + Y|\mathcal{G}] = \mathbb{E}[X|\mathcal{G}] + \mathbb{E}[Y|\mathcal{G}]$ almost surely.
- (4) If Y is $\mathcal{G}$ -measurable and $\mathbb{E}[XY] < \infty$, then $\mathbb{E}[XY|\mathcal{G}] = Y\mathbb{E}[X|\mathcal{G}]$ almost surely.
- (5) If $\mathcal{G}_1 \subseteq \mathcal{G}_2$, then $\mathbb{E}[X|\mathcal{G}_1] = \mathbb{E}[\mathbb{E}[X|\mathcal{G}_2]|\mathcal{G}_1]$ almost surely.

1.3 Stochastic Processes and Markov Chains

Definition. Let two measurable spaces $(\mathcal{X}, \mathcal{F})$ and $(\mathcal{Y}, \mathcal{G})$ be given.

- $(\mathcal{X}, \mathcal{F})$ and $(\mathcal{Y}, \mathcal{G})$ are called **isomorphic** if there exists a bijective function $f : \mathcal{X} \mapsto \mathcal{Y}$ such that $A \in \mathcal{F} \iff f(A) \in \mathcal{G}$.
- $(\mathcal{X}, \mathcal{F})$ is a **Borel space** if it is isomorphic to $(\mathbb{R}, \mathcal{B}(\mathbb{R}))$ for some set $A \in \mathcal{B}(\mathbb{R})$.

Theorem 3.1. Let μ be a probability measure on a Borel space S and λ be the Lebesgue measure on $([0, 1], \mathcal{B}([0, 1]))$. Then there exists a sequence of independent random element $X_1, X_2, \dots$ on $([0, 1], \mathcal{B}([0, 1]), \lambda)$ such that $\lambda_{X_t}(\cdot) := \lambda(X_t^{-1}(\cdot)) = \mu(\cdot)$ for all t .

Definition. For "countable" set $\mathcal{T}$, A stochastic process is a collection of random variables $X_t : t \in \mathcal{T}$.

Theorem 3.2. Let Borel spaces $(\Omega_n, \mathcal{F}_n)$ and measures μ_n on $(\Omega_1 \times \dots \times \Omega_n, \mathcal{F}_1 \otimes \dots \otimes \mathcal{F}_n)$ are given for each $n \in \mathbb{N}$.

$$\begin{aligned} &\forall n \in \mathbb{N} \text{ and } A \in \mathcal{F}_1 \otimes \dots \otimes \mathcal{F}_n, \quad \mu_{n+1}(A \times \Omega_{n+1}) = \mu_n(A) \\ &\implies \exists \text{ a probability space } (\Omega, \mathcal{F}, \mathbb{P}) \\ &\text{and random elements } X_1, X_2, \dots \text{ with } X_t : \Omega \rightarrow \Omega_t \text{ such that } \mathbb{P}_{X_1, \dots, X_n} = \mu_n \text{ for all } n. \end{aligned}$$

Definition. Let $(\mathcal{X}, \mathcal{F})$ and $(\mathcal{Y}, \mathcal{G})$ be measurable spaces. A **probability kernel** or **Markov kernel** from $(\mathcal{X}, \mathcal{F})$ to $(\mathcal{Y}, \mathcal{G})$ is a function $K : \mathcal{X} \times \mathcal{G} \rightarrow [0, 1]$ such that

- (a) $K(x, \cdot)$ is a measure for all $x \in \mathcal{X}$ and
- (b) $K(\cdot, A)$ is $\mathcal{F}$ -measurable for all $A \in \mathcal{G}$.

Notation: $K_x(A)$ or $K(A|x)$

Definition. Let K_1 be a probability kernel from $(\mathcal{X}, \mathcal{F}) \rightarrow (\mathcal{Y}, \mathcal{G})$ and K_2 be a probability kernel from $(\mathcal{Y}, \mathcal{G}) \rightarrow (\mathcal{Z}, \mathcal{H})$. Then the product kernel $K_1 \otimes K_2$ is the probability kernel from $(\mathcal{X}, \mathcal{F})$ to $(\mathcal{Y} \times \mathcal{Z}, \mathcal{G} \otimes \mathcal{H})$ defined by

$$(K_1 \otimes K_2)(x, A) = \int_{\mathcal{Y}} \int_{\mathcal{Z}} \mathbb{1}_A((y, z)) K_2(y, dz) K_1(x, dy).$$

Theorem 3.3. Let $(\Omega_n, \mathcal{F}_n)_{n=1}^{\infty}$ be a sequence of measurable spaces.

K_1 : a probability measure on $(\Omega_1, \mathcal{F}_1)$,

K_n : a probability kernel from $\Pi_{t=1}^{n-1} \Omega_t$ to Ω_n for $n \geq 2$.

$\implies \exists$ a probability space $(\Omega, \mathcal{F}, \mathbb{P})$

and random elements $(X_t)_{t=1}^{\infty}$ with $X_t : \Omega \rightarrow \Omega_t$ such that $\mathbb{P}_{X_1, \dots, X_n} = \bigotimes_{t=1}^n K_t$ for all n .

Definition. A **homogeneous Markov chain** is a sequence of random elements $(X_t)_{t=1}^{\infty}$ taking values in state space $\mathcal{S} = (\mathcal{X}, \mathcal{F})$ and with

$$\mathbb{P}(X_{t+1} \in \cdot | X_1, \dots, X_t) = \mathbb{P}(X_{t+1} \in \cdot | X_t) = \mu(X_t, \cdot), \quad \text{a.s.}$$

Here μ is a probability kernel from $(\mathcal{X}, \mathcal{F})$ to $(\mathcal{X}, \mathcal{F})$.

Definition 3.4. A $\mathbb{F}$ -adapted sequence of random variables $(X_t)_{t \in \mathbb{N}_+}$ (i.e. X_t is $\mathcal{F}_t$ -measurable for all $t \in \mathbb{N}$) is a $\mathbb{F}$ -adapted **martingale** if

(a) $\mathbb{E}[|X_t|] < \infty$ for all t .

(b) $\mathbb{E}[X_{t+1} | \mathcal{F}_t] = X_t$ for all t .

Supermartingale if less-than, submartingale if greater-than.

Example 3.5. $(Y_t)_{t \in \mathbb{N}}$ be a sequence of independent random variables with $\mathbb{P}(Y_t = 1) = \mathbb{P}(Y_t = -1) = \frac{1}{2}$. Then the winnings after t rounds is $S_t = \sum_{i=1}^t Y_i$, a martingale adapted to the filtration $(\mathcal{F}_t)_{t \in \mathbb{N}_+}$, where $\mathcal{F}_t = \sigma(Y_1, \dots, Y_t)$.

Definition 3.6. Let $\mathbb{F} = (\mathcal{F}_t)_{t \in \mathbb{N}}$ be a filtration. A random variable $\tau \in \mathbb{N} \cup \{\infty\}$ is a **stopping time** with respect to the filtration $\mathbb{F}$ if

$$\{\tau \leq t\} \in \mathcal{F}_t \text{ for all } t \in \mathbb{N}.$$

The σ -algebra generated by the stopping time τ is

$$\mathcal{F}_\tau = \{A \in \mathcal{F}_\infty : A \cap \{\tau \leq t\} \in \mathcal{F}_t \text{ for all } t \in \mathbb{N}\}.$$

Theorem 3.8 (Doob's Optional Stopping Theorem). Let $\mathbb{F} = (\mathcal{F}_t)_{t \in \mathbb{N}}$ be a filtration and $(X_t)_{t \in \mathbb{N}}$ be an $\mathbb{F}$ -adapted martingale. Let τ be a stopping time with respect to $\mathbb{F}$, and one of the following conditions is satisfied:

(a) There exists an $n \in \mathbb{N}$ such that $\mathbb{P}(\tau > n) = 0$.

(b) $\mathbb{E}[\tau] < \infty$ and there exists a constant $c \in \mathbb{R}$ such that $|X_{t+1} - X_t| \leq c$ for all t , $\mathbb{E}[|X_{t+1} - X_t| | \mathcal{F}_t] \leq c$ a.s. for $\tau > t$.

(c) There exists a constant c such that $|X_{t \wedge \tau}| \leq c$ a.s. for all $t \in \mathbb{N}$.

Then X_τ is almost surely well-defined and

$$\mathbb{E}[X_\tau] = \mathbb{E}[X_0].$$

Theorem 3.9 (Maximal inequality). Let $(X_t)_{t=0}^{\infty}$ be a supermartingale with $X_t \geq 0$ a.s. for all $t \in \mathbb{N}$. Then for any $\epsilon > 0$,

$$\mathbb{P}\left(\sup_{t \in \mathbb{N}} X_t \geq \epsilon\right) \leq \frac{\mathbb{E}[X_0]}{\epsilon}.$$

Theorem 3.10. Let $(X_t)_{t=0}^{\infty}$ be a submartingale with $X_t \geq 0$ a.s. for all $t \in \mathbb{N}$. Then for any $\epsilon > 0$,

$$\mathbb{P}\left(\max_{t \in 0, 1, \dots, n} X_t \geq \epsilon\right) \leq \frac{\mathbb{E}[X_n]}{\epsilon}.$$

1.4 Stochastic Bandits

Definition. A **stochastic bandit** is a collection of distributions $\nu = \{P_a : a \in \mathcal{A}\}$.

Here we are challenging to maximize the total reward $S_n = \sum_{t=1}^n X_t$ over n rounds.

Definition. A set of bandits $\mathcal{E}$ for which $\nu \in \mathcal{E}$ is called the environment class.

- An environment class $\mathcal{E}$ is **unstructured** if $\mathcal{A}$ is finite and

$$\forall a \in \mathcal{A}, \exists \text{ sets of distributions } \mathcal{M}_a \text{ such that } \mathcal{E} = \{\nu = (P_a : a \in \mathcal{A}) : P_a \in \mathcal{M}_a, \forall a \in \mathcal{A}\}.$$

or, in short, $\mathcal{E} = \times_{a \in \mathcal{A}} \mathcal{M}_a$.

- Environment classes with finite number of degrees of freedom are called **parametric** (Bernoulli).
- Environment classes that are not parametric are called **nonparametric** (subgaussian).
- An environment class is called **structured** if it is not unstructured.

Example 4.1. Let $\mathcal{A} = \{1, 2\}$ and $\mathcal{E} = \{(\mathcal{B}(\theta), \mathcal{B}(1 - \theta)) : \theta \in [0, 1]\}$ where $\mathcal{B}(\theta)$ denotes the Bernoulli distribution with mean θ . Then $\mathcal{E}$ is a structured parametric environment class.

Example 4.2 (Stochastic linear bandit). Let $\mathcal{A} \subseteq \mathbb{R}^d$ and $\theta \in \mathbb{R}^d$ and

$$\nu_\theta = (\mathcal{N}(\langle a, \theta \rangle, 1) : a \in \mathcal{A}) \quad \text{and} \quad \mathcal{E} = \{\nu_\theta : \theta \in \mathbb{R}^d\}.$$

Example 4.3. Consider an undirected graph G with vertices $V = 1, \dots, |V|$ and edges $E = 1, \dots, |E|$. In each round the learner chooses a path from vertex 1 to vertex $|V|$. Then each edge $e \in [E]$ is removed from the graph with probability $1 - \theta_e$ for unknown $\theta \in [0, 1]^{|E|}$. In this problem, $\mathcal{A}$ is the set of paths and

$$\nu_\theta = \left(\mathcal{B} \left(\prod_{e \in a} \theta_e \right) : a \in \mathcal{A} \right) \quad \text{and} \quad \mathcal{E} = \{\nu_\theta : \theta \in [0, 1]^{|E|}\}.$$

Definition. Let $\nu = (P_a : a \in \mathcal{A})$ be a stochastic bandit.

$$\mu_a(\nu) := \int_{-\infty}^{\infty} x dP_a(x)$$

Also let $\mu^*(\nu) := \max_{a \in \mathcal{A}} \mu_a(\nu)$, the largest mean of all the arms.

Definition. The **regret** of policy π on bandit instance ν is

$$R_n(\pi, \nu) := n\mu^*(\nu) - \mathbb{E} \left[\sum_{t=1}^n X_t \right]$$

where the expectation is taken with respect to the probability measure on outcomes induced by the interaction of π and ν .

Lemma 4.4. Let ν be a stochastic bandit environment. Then

- $R_n(\pi, \nu) \geq 0$ for all policies π ,
- the policy π choosing $A_t \in \arg \max_a \mu_a(\nu)$ for all t satisfies $R_n(\pi, \nu) = 0$, and
- if $R_n(\pi, \nu) = 0$ for some policy π , then $\mathbb{P}[\mu_{A_t}(\nu) = \mu^*(\nu)] = 1$ for all $t \in [n]$.

Proof. (a) Since $\mu^*(\nu) = \max_{a \in \mathcal{A}} \mu_a(\nu)$,

$$n\mu^*(\nu) \geq \mathbb{E} \left[\sum_{t=1}^n \mu_{A_t}(\nu) \right].$$

By the linearity of expectation,

$$\mathbb{E} \left[\sum_{t=1}^n \mu_{A_t}(\nu) \right] = \sum_{t=1}^n \mathbb{E} [\mu_{A_t}(\nu)].$$

By the stochastic bandit assumption and the tower property,

$$\begin{aligned} \sum_{t=1}^n \mathbb{E} [\mu_{A_t}(\nu)] &= \sum_{t=1}^n \mathbb{E} [\mathbb{E} [X_t \mid H_{t-1}, A_t]] \\ &= \sum_{t=1}^n \mathbb{E} [X_t] \\ &= \mathbb{E} \left[\sum_{t=1}^n X_t \right], \end{aligned}$$

using linearity of expectation again. Therefore,

$$R_n(\pi, \nu) = n\mu^*(\nu) - \mathbb{E} \left[\sum_{t=1}^n X_t \right] \geq 0.$$

(b) Let π be the policy choosing $A_t \in \arg \max_a \mu_a(\nu)$ for all t . Then $\mu^*(\nu) = \mu_{A_t}(\nu)$ for all t , which implies

$$\begin{aligned} R_n(\pi, \nu) &= n\mu^*(\nu) - \mathbb{E} \left[\sum_{t=1}^n X_t \right] \\ &= \sum_{t=1}^n (\mu^*(\nu) - \mathbb{E} [\mu_{A_t}(\nu)]) \\ &= \mathbb{E} \left[\sum_{t=1}^n (\mu^*(\nu) - \mu_{A_t}(\nu)) \right] \\ &= 0 \end{aligned}$$

and therefore $R_n(\pi, \nu) = 0$.

(c) Suppose $R_n(\pi, \nu) = 0$ for some policy π . Then

$$\sum_{t=1}^n \mathbb{E} [\mu^*(\nu) - \mu_{A_t}(\nu)] = 0.$$

Since $\mu^*(\nu) - \mu_{A_t}(\nu) \geq 0$ for all t , $\mu^*(\nu) = \mu_{A_t}(\nu)$ almost surely for all $t \in \{1, \dots, n\}$. □

Remark (Factorisation of the regret). The learner cannot always achieve zero regret, so the relatively weak objective is to find a policy π with

$$\forall \nu \in \mathcal{E}, \quad \lim_{n \rightarrow \infty} \frac{R_n(\pi, \nu)}{n} = 0.$$

Or further objective, finding some constants $C > 0$ and $p < 1$ such that

$$\forall \nu \in \mathcal{E}, \quad R_n(\pi, \nu) \leq Cn^p.$$

Another alternative is to find a function $C : \mathcal{E} \rightarrow [0, \infty)$ and $f : \mathbb{N} \rightarrow [0, \infty)$ such that

$$\forall \nu \in \mathcal{E}, \quad R_n(\pi, \nu) \leq C(\nu)f(n).$$

Bayesian regret is another candidate objective. For prior probability measure Q on $\mathcal{E}$, the Bayesian regret is the average of the regret with respect to the prior Q , i.e.,

$$\text{BR}_n(\pi, Q) = \int_{\mathcal{E}} R_n(\pi, \nu) dQ(\nu)$$

where the regret R_n is measurable.

Definition. Let $\nu = (P_a : a \in \mathcal{A})$ be a stochastic bandit. Then the **suboptimality gap** of action a is defined by

$$\Delta_a(\nu) := \mu^*(\nu) - \mu_a(\nu).$$

Also let

$$T_a(t) := \sum_{s=1}^t \mathbb{1}\{A_s = a\}$$

be the number of times action a has been chosen up to time t .

Lemma 4.5 (Regret decomposition lemma). *For any policy π and stochastic bandit ν with $\mathcal{A}$ finite or countable and horizon $n \in \mathbb{N}$, the regret R_n of policy π in ν satisfies*

$$R_n(\pi, \nu) = \sum_{a \in \mathcal{A}} \Delta_a(\nu) \mathbb{E}[T_a(n)].$$

Proof. For any $t \in \mathbb{N}$, $\sum_{a \in \mathcal{A}} \mathbb{1}\{A_t = a\} = 1$. Hence $S_n = \sum_t X_t = \sum_t \sum_a X_t \mathbb{1}\{A_t = a\}$. That is

$$R_n = n\mu^* - \mathbb{E}[S_n] = \sum_{a \in \mathcal{A}} \sum_{t=1}^n \mathbb{E}[(\mu^* - X_t) \mathbb{1}\{A_t = a\}].$$

Since $\mathbb{E}[X_t \mid H_{t-1}, A_t] = \mu_{A_t}$, we have

$$\begin{aligned} \mathbb{E}[(\mu^* - X_t) \mathbb{1}\{A_t = a\} \mid A_t] &= \mathbb{1}\{A_t = a\} \mathbb{E}[\mu^* - X_t \mid A_t] \\ &= \mathbb{1}\{A_t = a\} (\mu^* - \mu_a) \\ &= \mathbb{1}\{A_t = a\} \Delta_a. \end{aligned}$$

Therefore

$$\begin{aligned} R_n &= \sum_{a \in \mathcal{A}} \sum_{t=1}^n \mathbb{E}[(\mu^* - X_t) \mathbb{1}\{A_t = a\}] \\ &= \sum_{a \in \mathcal{A}} \sum_{t=1}^n \mathbb{E}[\mathbb{E}[(\mu^* - X_t) \mathbb{1}\{A_t = a\} \mid A_t]] \\ &= \sum_{a \in \mathcal{A}} \sum_{t=1}^n \mathbb{E}[\Delta_a \mathbb{1}\{A_t = a\}] \\ &= \sum_{a \in \mathcal{A}} \Delta_a \mathbb{E}[T_a(n)]. \end{aligned}$$

□

Exercise 4.11(a) (Retirement policy). *Proof.* Let

$$\mathcal{H}_j := (\{1, 2\} \times \mathbb{R})^j, \quad j \in \{0, \dots, n\},$$

where $\mathcal{H}_0 = \{\emptyset\}$. Fix $n \in \mathbb{N}$ and a policy $\pi = (\pi_j)_{j=1}^n$. Let $U_1, \dots, U_n \stackrel{\text{i.i.d.}}{\sim} \text{Unif}(0, 1)$. For $j \in \{0, \dots, n-1\}$ and $h \in \mathcal{H}_j$, define

$$\psi_{j+1}(h) := \begin{cases} 1, & U_{j+1} \leq \pi_{j+1}(1 \mid h), \\ 2, & U_{j+1} > \pi_{j+1}(1 \mid h). \end{cases}$$

Define

$$\Gamma : \bigcup_{j=0}^n \mathcal{H}_j \longrightarrow \bigcup_{j=0}^n \mathcal{H}_j$$

such that $\forall j \in \{0, \dots, n\}$ and $\forall h \in \mathcal{H}_j$,

$$\Gamma(h) := \begin{cases} h, & j = n, \\ h, & j < n \text{ and } \psi_{j+1}(h) = 1, \\ \Gamma(h, (2, \mu_2)), & j < n \text{ and } \psi_{j+1}(h) = 2. \end{cases}$$

Now fix $\nu = (\nu_1, \delta_{\mu_2}) \in \mathcal{E}$. Let

$$Y_1, Y_2, \dots \stackrel{\text{i.i.d.}}{\sim} \nu_1$$

be the rewards obtained from arm 1. Define $\tilde{H}_0, \dots, \tilde{H}_n$ and $A'_1, \dots, A'_n$ by

$$\tilde{H}_0 := \Gamma(\emptyset),$$

and, for $t \in \{1, \dots, n\}$,

$$A'_t := \begin{cases} 1, & |\tilde{H}_{t-1}| < n, \\ 2, & |\tilde{H}_{t-1}| = n, \end{cases}$$

and

$$\tilde{H}_t := \begin{cases} \Gamma(\tilde{H}_{t-1}, (1, Y_t)), & |\tilde{H}_{t-1}| < n, \\ \tilde{H}_{t-1}, & |\tilde{H}_{t-1}| = n. \end{cases}$$

Then $\forall t \in \{0, \dots, n\}$, $\tilde{H}_t \in \bigcup_{j=0}^n \mathcal{H}_j$. Define $\forall h \in \mathcal{H}_{t-1}$ and $\forall a \in \{1, 2\}$,

$$H'_{t-1} \in \mathcal{H}_{t-1} \text{ and } \pi'_t(a \mid h) = \mathbb{P}(A'_t = a \mid H'_{t-1} = h).$$

for all $t \in \{1, \dots, n\}$. We first show that π' is a retirement policy. Suppose that $A'_{t_0} = 2$ for some $t_0 \in \{1, \dots, n\}$. Then

$$|\tilde{H}_{t_0-1}| = n.$$

Hence, by the definition of $\tilde{H}_t$,

$$\tilde{H}_s = \tilde{H}_{t_0-1} \quad \text{for all } s \geq t_0 - 1.$$

Therefore,

$$A'_s = 2 \quad \text{for all } s \geq t_0.$$

Thus, π' is a retirement policy. Let

$$K := T_1^{\pi'}(n) = \sum_{t=1}^n \mathbb{1}_{\{A'_t=1\}}.$$

By construction,

$$\tilde{H}_K \in \mathcal{H}_n.$$

Moreover, conditionally on each prefix of $\tilde{H}_K$, its next action is chosen according to the same conditional distribution as under π , and its next reward has distribution ν_1 if that action is 1, and equals μ_2 if that action is 2. Consequently,

$$\tilde{H}_K \stackrel{d}{=} H_n^\pi,$$

where H_n^π denotes the n -round history generated by π in the bandit ν . Since K is exactly the number of occurrences of action 1 in $\tilde{H}_K$, it follows that

$$T_1^{\pi'}(n) \stackrel{d}{=} T_1^\pi(n).$$

Hence,

$$\mathbb{E}_\nu^{\pi'}[T_1^{\pi'}(n)] = \mathbb{E}_\nu^\pi[T_1^\pi(n)].$$

Since $T_2(n) = n - T_1(n)$,

$$\mathbb{E}_\nu^{\pi'}[T_2^{\pi'}(n)] = \mathbb{E}_\nu^\pi[T_2^\pi(n)].$$

Finally, by the regret decomposition lemma,

$$\begin{aligned} R_n(\pi', \nu) &= \Delta_1(\nu) \mathbb{E}_\nu^{\pi'}[T_1^{\pi'}(n)] + \Delta_2(\nu) \mathbb{E}_\nu^{\pi'}[T_2^{\pi'}(n)] \\ &= \Delta_1(\nu) \mathbb{E}_\nu^\pi[T_1^\pi(n)] + \Delta_2(\nu) \mathbb{E}_\nu^\pi[T_2^\pi(n)] \\ &= R_n(\pi, \nu). \end{aligned}$$

Therefore,

$$R_n(\pi', \nu) \leq R_n(\pi, \nu)$$

for every $\nu \in \mathcal{E}$. □

1.5 Concentration of measure

Definition. Let $X, X_1, X_2, \dots, X_n$ be i.i.d. random variables with real values, and assume that the mean $\mu = \mathbb{E}[X]$ and variance $\sigma^2 = \mathbb{V}[X] < \infty$. Then the **empirical mean** is defined by

$$\hat{\mu} := \frac{1}{n} \sum_{i=1}^n X_i.$$

This is an unbiased estimator of the mean μ , since $\mathbb{E}[\hat{\mu}] = \mathbb{E}[\frac{1}{n} \sum_{i=1}^n X_i] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \mu$. Using independence of the X_i 's, the variance of the empirical mean is

$$\mathbb{V}[\hat{\mu}] = \mathbb{E}[(\hat{\mu} - \mu)^2] = \frac{\sigma^2}{n}.$$

Two quantities

$$\mathbb{P}[\hat{\mu} \geq \mu + \epsilon] \quad \text{and} \quad \mathbb{P}[\hat{\mu} \leq \mu - \epsilon]$$

are called the **tail probabilities** of $\hat{\mu} - \mu$.

Lemma 5.1. *Let X be a random variable and let $\epsilon > 0$. Then*

$$(a) \text{ (Markov's inequality) : } \mathbb{P}[|X| \geq \epsilon] \leq \frac{\mathbb{E}[|X|]}{\epsilon}.$$

$$(b) \text{ (Chebyshev's inequality) : } \mathbb{P}[|X - \mu| \geq \epsilon] \leq \frac{\sigma^2}{\epsilon^2}.$$

Proof. (a)

$$\begin{aligned} \mathbb{E}[|X|] &= \int_0^\infty \mathbb{P}[|X| \geq t] dt \\ &\geq \int_0^\epsilon \mathbb{P}[|X| \geq \epsilon] dt \\ &= \epsilon \mathbb{P}[|X| \geq \epsilon]. \end{aligned}$$

(b)

$$\begin{aligned} \mathbb{P}[|X - \mu| \geq \epsilon] &= \mathbb{P}[(X - \mu)^2 \geq \epsilon^2] \\ &\leq \frac{\mathbb{E}[(X - \mu)^2]}{\epsilon^2}. \end{aligned}$$

□

Theorem (Central Limit Theorem). Let $X_1, X_2, \dots, X_n$ be i.i.d. random variables with finite second moment. Define

$$S_n := \sum_{t=1}^n (X_t - \mu).$$

Then for $\sigma^2 > 0$,

$$\frac{S_n}{\sqrt{n}\sigma} \xrightarrow{d} N(0, 1).$$

Definition 5.2. A random variable X is σ -subgaussian if

$$\forall \lambda \in \mathbb{R}, \quad \mathbb{E}[e^{\lambda X}] \leq e^{\frac{\lambda^2 \sigma^2}{2}}.$$

The **moment generating function** of X is a function

$$M_X(\lambda) := \mathbb{E}[e^{\lambda X}], \quad \lambda \in \mathbb{R}.$$

The **cumulant generating function** of X is a function

$$\Lambda_X(\lambda) := \log M_X(\lambda) = \log \mathbb{E}[e^{\lambda X}], \quad \lambda \in \mathbb{R}.$$

Remark. X is σ -subgaussian if and only if

$$\Lambda_X(\lambda) \leq \frac{\lambda^2 \sigma^2}{2} \quad \forall \lambda \in \mathbb{R}.$$

Theorem 5.3. If X is σ -subgaussian, then for all $\epsilon \geq 0$,

$$\mathbb{P}[X \geq \epsilon] \leq \exp\left(-\frac{\epsilon^2}{2\sigma^2}\right).$$

Proof. Let $\lambda > 0$. By Markov's inequality,

$$\mathbb{P}[X \geq \epsilon] = \mathbb{P}[e^{\lambda X} \geq e^{\lambda \epsilon}] \leq \frac{\mathbb{E}[e^{\lambda X}]}{e^{\lambda \epsilon}} \leq \frac{e^{\frac{\lambda^2 \sigma^2}{2}}}{e^{\lambda \epsilon}} = \exp\left(\frac{\lambda^2 \sigma^2}{2} - \lambda \epsilon\right).$$

Letting $\lambda = \frac{\epsilon}{\sigma^2}$ gives

$$\mathbb{P}[X \geq \epsilon] \leq \exp\left(-\frac{\epsilon^2}{2\sigma^2}\right).$$

□

Lemma 5.4. Suppose X is σ -subgaussian and X_1, X_2 are independent σ_1, σ_2 -subgaussian, respectively. Then

- (a) $\mathbb{E}[X] = 0$ and $\mathbb{V}[X] \leq \sigma^2$.
- (b) cX is $|c|\sigma$ -subgaussian for all $c \in \mathbb{R}$.
- (c) $X_1 + X_2$ is $\sqrt{\sigma_1^2 + \sigma_2^2}$ -subgaussian.

Proof. $\log x$ is convex on $(0, \infty)$, in the sense that for any $x, y > 0$ and $\theta \in [0, 1]$,

$$\log(\theta x + (1 - \theta)y) \geq \theta \log(x) + (1 - \theta) \log(y).$$

- (a) By Jensen's inequality, $\log e^{\frac{\lambda^2 \sigma^2}{2}} \geq \log \mathbb{E}[e^{\lambda X}] \geq \mathbb{E}[\log e^{\lambda X}] = \lambda \mathbb{E}[X]$. That is, $\mathbb{E}[X] \leq \frac{\lambda \sigma^2}{2}$ for all $\lambda \in \mathbb{R}$, which implies $\mathbb{E}[X] = 0$. Also,

$$\mathbb{E}[e^{\lambda X}] + \mathbb{E}[e^{-\lambda X}] \leq 2e^{\frac{\lambda^2 \sigma^2}{2}}, \quad \forall \lambda \in \mathbb{R}.$$

Hence,

$$\mathbb{E} \left[\frac{e^{\lambda X} - e^{-\lambda X} - 2}{\lambda^2} \right] \leq \frac{e^{\frac{2(\lambda^2 \sigma^2) - 1}{2}}}{\lambda^2}.$$

Here

$$\frac{e^{\lambda X} - e^{-\lambda X}}{\lambda^2} = \frac{2(\cosh(\lambda X) - 1)}{\lambda^2},$$

which is non-negative and $\frac{2(\cosh(\lambda X) - 1)}{\lambda^2} \rightarrow X^2$ as $\lambda \rightarrow 0$. Therefore, by Fatou's lemma,

$$\mathbb{E}[X^2] \leq \liminf_{\lambda \rightarrow 0} \mathbb{E} \left[\frac{2(\cosh(\lambda X) - 1)}{\lambda^2} \right] \leq \lim_{\lambda \rightarrow 0} \frac{e^{\frac{2(\lambda^2 \sigma^2) - 1}{2}}}{\lambda^2} = \sigma^2.$$

(b) $\forall \lambda \in \mathbb{R}, \mathbb{E}[e^{\lambda cX}] = \mathbb{E}[e^{(c\lambda)X}] \leq e^{\frac{(c\lambda)^2 \sigma^2}{2}} = e^{\frac{\lambda^2 (c\sigma)^2}{2}}$, which shows that cX is $|c|\sigma$ -subgaussian.

(c) $\forall \lambda \in \mathbb{R}, \mathbb{E}[e^{\lambda(X_1 + X_2)}] = \mathbb{E}[e^{\lambda X_1}] \mathbb{E}[e^{\lambda X_2}] \leq e^{\frac{\lambda^2 \sigma_1^2}{2}} e^{\frac{\lambda^2 \sigma_2^2}{2}} = e^{\frac{\lambda^2 (\sigma_1^2 + \sigma_2^2)}{2}}$, which shows that $X_1 + X_2$ is $\sqrt{\sigma_1^2 + \sigma_2^2}$ -subgaussian. □

Corollary 5.5. Assume $X_i - \mu$ are independent, σ -subgaussian random variables. Then for any $\epsilon \geq 0$,

$$\mathbb{P}(\hat{\mu} \geq \mu + \epsilon) \leq \exp \left(-\frac{n\epsilon^2}{2\sigma^2} \right) \quad \text{and} \quad \mathbb{P}(\hat{\mu} \leq \mu - \epsilon) \leq \exp \left(-\frac{n\epsilon^2}{2\sigma^2} \right),$$

where $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$.

Proof. By Lemma 5.4, $\hat{\mu} - \mu$ is $\frac{\sigma}{\sqrt{n}}$ -subgaussian. Therefore, for any $\epsilon \geq 0$,

$$\mathbb{P}(\hat{\mu} \geq \mu + \epsilon) \leq \exp \left(-\frac{n\epsilon^2}{2\sigma^2} \right) \quad \text{and} \quad \mathbb{P}(\hat{\mu} \leq \mu - \epsilon) \leq \exp \left(-\frac{n\epsilon^2}{2\sigma^2} \right)$$

by Theorem 5.3. □

1.6 The Explore-Then-Commit Algorithm

From now on, we assume that all bandit instances are in $\mathcal{E}_{\text{SG}}^k(1)$, which means the reward distribution of all arms is 1-subgaussian. Let $\hat{\mu}_i(t)$ be the average reward received from arm i up to round t . Formally,

$$\hat{\mu}_i(t) = \frac{1}{T_i(t)} \sum_{s=1}^t \mathbb{1}\{A_s = i\} X_s,$$

where $T_i(t) = \sum_{s=1}^t \mathbb{1}\{A_s = i\}$ is the number of times arm i has been played up to round t .

Input: m

1 In round t choose action

$$A_t = \begin{cases} (t \bmod k) + 1, & \text{if } t \leq mk, \\ \operatorname{argmax}_i \hat{\mu}_i(mk) & \text{if } t > mk, \end{cases}$$

Algorithm 1: ETC Algorithm

Theorem 6.1. Suppose ETC with 1-subgaussian bandit and $1 \leq m \leq n/k$. Then

$$R_n \leq m \sum_{i=1}^k \Delta_i + (n - mk) \sum_{i=1}^k \Delta_i \exp \left(-\frac{m\Delta_i^2}{4} \right)$$

Proof. □

1.7 The Upper Confidence Bound Algorithm

Note that for sequence of independent 1-subgaussian random variables $(X_t)_{t=1}^n$, mean μ and the empirical mean $\hat{\mu}_n = \frac{1}{n} \sum_{t=1}^n X_t$ satisfy

$$\mathbb{P} \left(\mu \geq \hat{\mu}_n + \sqrt{\frac{2 \log(1/\delta)}{n}} \right) \leq \delta, \quad \forall \delta \in (0, 1).$$

Define

$$\text{UCB}_i(t-1, \delta) = \begin{cases} \infty, & \text{if } T_i(t-1) = 0, \\ \hat{\mu}_i(t-1) + \sqrt{\frac{2 \log(1/\delta)}{T_i(t-1)}}, & \text{if } T_i(t-1) > 0. \end{cases}$$

1 Input k and δ
2 for $t \in \{1, \dots, n\}$ do
3 Choose action

$$A_t = \underset{i}{\operatorname{argmax}} \text{UCB}_i(t-1, \delta)$$

4 Observe reward X_t and update upper confidence bounds
5 end for

Algorithm 2: UCB(δ)

Theorem 7.1. Suppose UCB with stochastic k -armed 1-subgaussian bandits. Then for all n , if $\delta = 1/n$, then

$$R_n \leq 3 \sum_{i=1}^k \Delta_i + \sum_{i: \Delta_i > 0} \frac{16 \log(n)}{\Delta_i}$$

Proof. WLOG, assume that $\mu_1 = \mu^*$. Let G_i be the 'good' event defined by

$$G_i = \left\{ \mu_1 < \min_{t \in \{1, \dots, n\}} \text{UCB}_1(t, \delta) \right\} \cap \left\{ \hat{\mu}_{i, u_i} + \sqrt{\frac{2}{u_i} \log \left(\frac{1}{\delta} \right)} < \mu_1 \right\}$$

where $u_i \in \{1, \dots, n\}$. Assume that G_i holds for all $i \in \{1, \dots, k\}$. To make contradiction, suppose that $T_i(n) > u_i$. Then there exists $t \in \{1, \dots, n\}$ such that $T_i(t-1) = u_i$ and $A_t = i$. Hence

$$\begin{aligned} \text{UCB}_i(t-1, \delta) &= \hat{\mu}_i(t-1) + \sqrt{\frac{2 \log(1/\delta)}{T_i(t-1)}} \\ &= \hat{\mu}_{i, u_i} + \sqrt{\frac{2}{u_i} \log \left(\frac{1}{\delta} \right)} \\ &< \mu_1 \\ &\leq \text{UCB}_1(t-1, \delta). \end{aligned}$$

Therefore, $A_t = \operatorname{argmax}_j \text{UCB}_j(t-1, \delta) \neq i$, which is a contradiction. On the other hand,

$$G_i^c = \left\{ \mu_1 \geq \min_{t \in \{1, \dots, n\}} \text{UCB}_1(t, \delta) \right\} \cup \left\{ \hat{\mu}_{i, u_i} + \sqrt{\frac{2}{u_i} \log \left(\frac{1}{\delta} \right)} \geq \mu_1 \right\}.$$

Former part bounded by

$$\begin{aligned}
\mathbb{P}\left(\mu_1 \geq \min_{t \in \{1, \dots, n\}} \text{UCB}_1(t, \delta)\right) &\leq \mathbb{P}\left(\mu_1 \geq \min_{s \in \{1, \dots, n\}} \hat{\mu}_1(s) + \sqrt{\frac{2}{s} \log\left(\frac{1}{\delta}\right)}\right) \\
&\leq \mathbb{P}\left(\bigcup_{s \in \{1, \dots, n\}} \left\{\mu_1 \geq \hat{\mu}_1(s) + \sqrt{\frac{2}{s} \log\left(\frac{1}{\delta}\right)}\right\}\right) \\
&\leq \sum_{s \in \{1, \dots, n\}} \mathbb{P}\left(\mu_1 \geq \hat{\mu}_1(s) + \sqrt{\frac{2}{s} \log\left(\frac{1}{\delta}\right)}\right) \\
&\leq \sum_{s \in \{1, \dots, n\}} \delta = n\delta.
\end{aligned}$$

For latter part, assume u_i is chosen large enough that

$$\Delta_i - \sqrt{\frac{2}{u_i} \log\left(\frac{1}{\delta}\right)} \geq c\Delta_i$$

for some $c \in (0, 1)$. Then

$$\begin{aligned}
\mathbb{P}\left(\hat{\mu}_{i, u_i} + \sqrt{\frac{2}{u_i} \log\left(\frac{1}{\delta}\right)} \geq \mu_1\right) &= \mathbb{P}\left(\hat{\mu}_{i, u_i} - \mu_i \geq \Delta_i - \sqrt{\frac{2}{u_i} \log\left(\frac{1}{\delta}\right)}\right) \\
&\leq \mathbb{P}(\hat{\mu}_{i, u_i} - \mu_i \geq c\Delta_i) \\
&\leq \exp\left(-\frac{u_i(c\Delta_i)^2}{2}\right)
\end{aligned}$$

by Cor.5.5. Now we have

$$\mathbb{P}(G_i^c) \leq n\delta + \exp\left(-\frac{u_i(c\Delta_i)^2}{2}\right).$$

Also

$$\mathbb{E}[T_i(n)] = \mathbb{E}[\mathbb{1}_{G_i} T_i(n)] + \mathbb{E}[\mathbb{1}_{G_i^c} T_i(n)].$$

Therefore, choosing

$$u_i = \left\lceil \frac{2 \log(1/\delta)}{((1-c)\Delta_i)^2} \right\rceil$$

follows that

$$\mathbb{E}[T_i(n)] \leq u_i + 1 + n^{1-2c^2/(1-c)^2} = \left\lceil \frac{2 \log(n^2)}{((1-c)\Delta_i)^2} \right\rceil + 1 + n^{1-2c^2/(1-c)^2}.$$

Lastly, choosing $c = 1/2$ follows that

$$\mathbb{E}[T_i(n)] \leq \left\lceil \frac{16 \log(n)}{\Delta_i^2} \right\rceil + 2,$$

and regret decomposition lemma leads to

$$\begin{aligned}
R(n) &\leq \sum_{i: \Delta_i > 0} \Delta_i \left(\left\lceil \frac{16 \log(n)}{\Delta_i^2} \right\rceil + 2 \right) \\
&= \sum_{i: \Delta_i > 0} \left(\frac{16 \log(n)}{\Delta_i} + \Delta_i \right) + 2 \sum_{i: \Delta_i > 0} \Delta_i \\
&\leq 3 \sum_{i=1}^k \Delta_i + \sum_{i: \Delta_i > 0} \frac{16 \log(n)}{\Delta_i}.
\end{aligned}$$

□

Theorem 7.2. Suppose $\nu \in \mathcal{E}_{\text{SG}}^k(1)$ and $\delta = 1/n^2$. Then regret of the UCB algorithm is bounded by

$$R(n) \leq 8\sqrt{nk \log(n)} + 3 \sum_{i=1}^k \Delta_i.$$

Proof. Let $\Delta > 0$ be some value to be determined later. Recall from proof of 7.1,

$$\mathbb{E}[T_i(n)] \leq 3 + \frac{16 \log(n)}{\Delta_i^2}.$$

Hence regret decomposition lemma follows that

$$\begin{aligned} R_n &= \sum_{i: \Delta_i > 0} \Delta_i \mathbb{E}[T_i(n)] \\ &= \sum_{i: \Delta_i < \Delta} \Delta_i \mathbb{E}[T_i(n)] + \sum_{i: \Delta_i \geq \Delta} \Delta_i \mathbb{E}[T_i(n)] \\ &\leq n\Delta + \sum_{i: \Delta_i \geq \Delta} \left(3\Delta_i + \frac{16 \log(n)}{\Delta_i} \right) \\ &\leq n\Delta + 3 \sum_{i=1}^k \Delta_i + \frac{16k \log(n)}{\Delta} \\ &\leq 8\sqrt{nk \log(n)} + 3 \sum_{i=1}^k \Delta_i. \end{aligned}$$

by choosing $\Delta = \sqrt{\frac{16k \log(n)}{n}}$. □

Exercise 7.4 (Phased UCB 1). Consider the following algorithm, called phased UCB;

```

1 Input  $k$  and  $\delta$ 
2 Choose each arm once
3 for  $l = 1, 2, \dots$  do
4   Compute  $A_l = \operatorname{argmax}_i \text{UCB}_i(t-1, \delta)$ 
5   Choose arm  $A_l$  and exactly  $2^l$  times
6 end for

```

Algorithm 3: A phased version of UCB

- (a) State and prove a bound on the regret for this version of UCB.
- (b) Compare with Theorem 7.1.
- (c) How would the result change if the l th phase had a length of $\lceil \alpha^l \rceil$ with $\alpha > 1$?

Solution.

- (a) Let $k = 2$, L be some large constant (≥ 3), $n = 2^{L+1}$, $\delta = n^{-2}$, and $q = 2\sqrt{\log(n)}$. Also let two arms attain deterministic reward, $\mu_1 = 0$, $\mu_2 = -\Delta$, respectively. ($\Delta > 0$) In this setting, the algorithm chooses each arm once initially and then goes to round 3. Here learner gets UCB

$$\text{UCB}_1(2, \delta) = q, \quad \text{UCB}_2(2, \delta) = -\Delta + q$$

and hence chooses arm 1 in round 3, and 4. Explicit simulation will be following;

t	l	UCB ₁	UCB ₂	Next action
1	0	∞	∞	1
2	0	q	∞	2
3	1	q	$-\Delta + q$	1
4	1	$q/\sqrt{2}$	$-\Delta + q$	1
5	2	$q/\sqrt{3}$	$-\Delta + q$	1
...	...	...	...	...
2^L	$L - 1$	$q/\sqrt{2^L - 2}$	$-\Delta + q$	1
$2^L + 1$	L	$q/\sqrt{2^L - 1}$	$-\Delta + q$	2
...	...	...	...	...

if we choose Δ with

$$\frac{q}{\sqrt{2^L - 1}} < q - \Delta < \frac{q}{\sqrt{2^{L-1} - 1}}.$$

Then the regret will be dominated by the suboptimal arm 2 in the last phase, giving

$$R_n \geq \Delta \cdot 2^L = \frac{n}{2} \Delta.$$

This example shows that the phased UCB algorithm can obtain linear regret,

$$R_n \leq n \cdot \max_i \Delta_i,$$

which is trivial.

(b) The above example of (a) shows that the phased UCB algorithm can get linear regret, while Theorem 7.1 shows that the standard UCB always get lower than sublinear regret.

(c) Still get linear regret for the phased UCB algorithm with any constant $\alpha > 1$.

Exercise 7.5 (Phased UCB 2). Consider the following another phased version of UCB;

```

1 Input  $k$  and  $\delta$ 
2 Choose each arm once
3 for  $l = 1, 2, \dots$  do
4   Let  $t_l = t$ 
5   Compute  $A_l = \operatorname{argmax}_i \text{UCB}_i(t_l - 1, \delta)$ 
6   Choose arm  $A_l$  until round  $t$  such that  $T_{A_l}(t) \geq \alpha T_{A_l}(t_l - 1)$ 
7 end for
```

(a) State and prove a bound on the regret for this version of UCB with $\alpha = 2$.

(b) Compare with previous exercise and Theorem 7.1.

(c) Repeat (a) and (b) for $\alpha > 1$.

(d) Implement theses algorithms and compare them empirically to UCB(δ)

Solution. WLOG, assume that arm 1 is optimal.

(a) For all suboptimal arms i , let

$$H_i = \left\{ \forall t \in \{1, \dots, n\}, \mu_1 < \hat{\mu}_{1,t} + \sqrt{\frac{2 \log(1/\delta)}{t}} \right\} \cap \left\{ \forall t \in \{1, \dots, n\}, \mu_i > \hat{\mu}_{i,t} - \sqrt{\frac{2 \log(1/\delta)}{t}} \right\}.$$

Then the complement event H_i^c is

$$H_i^c = \left\{ \exists t \in \{1, \dots, n\} : \mu_1 \geq \hat{\mu}_{1,t} + \sqrt{\frac{2 \log(1/\delta)}{t}} \right\} \cup \left\{ \exists t \in \{1, \dots, n\} : \mu_i \leq \hat{\mu}_{i,t} - \sqrt{\frac{2 \log(1/\delta)}{t}} \right\}.$$

That is

$$\mathbb{P}(H_i^c) \leq \sum_{t=1}^n \mathbb{P} \left(\mu_1 \geq \hat{\mu}_{1,t} + \sqrt{\frac{2 \log(1/\delta)}{t}} \right) + \sum_{t=1}^n \mathbb{P} \left(\mu_i \leq \hat{\mu}_{i,t} - \sqrt{\frac{2 \log(1/\delta)}{t}} \right).$$

Cor 5.5 follows that

$$\mathbb{P}(H_i^c) \leq 2 \sum_{t=1}^n \exp \left(-\frac{2 \log(1/\delta)}{t} \cdot \frac{t}{2} \right) = 2 \sum_{t=1}^n \delta = 2n\delta.$$

For the latest phase l where arm i is chosen, let $m := T_i(t_l - 1)$. Upon the event H_i , letting $m \geq u_i$ implies that

$$\text{UCB}_i(t_l - 1, \delta) = \hat{\mu}_{i,m} + \sqrt{\frac{2 \log(1/\delta)}{m}} < \mu_i + \sqrt{\frac{8 \log(1/\delta)}{m}} < \mu_i + \Delta_i = \mu_1 < \text{UCB}_1(t_l - 1, \delta)$$

where $u_i = \lfloor \frac{8 \log(1/\delta)}{\Delta_i^2} \rfloor + 1$.

On the other hand, if $m < u_i$, then $T_i(n) \leq 2m < 2u_i$. Therefore

$$\begin{aligned} \mathbb{E}[T_i(n)] &\leq 2u_i + n\mathbb{P}(H_i^c) \\ &= 2u_i + 2n^2\delta. \end{aligned}$$

Choosing $\delta = n^{-2}$ follows that $u_i \leq \frac{16 \log(n)}{\Delta_i^2} + 1$, and hence

$$\mathbb{E}[T_i(n)] \leq 2 \left(\frac{16 \log(n)}{\Delta_i^2} + 1 \right) + 2 = \frac{32 \log(n)}{\Delta_i^2} + 4.$$

The regret is bounded by

$$\begin{aligned} R(n) &= \sum_{i: \Delta_i > 0} \Delta_i \mathbb{E}[T_i(n)] \\ &\leq \sum_{i: \Delta_i > 0} \Delta_i \left(\frac{32 \log(n)}{\Delta_i^2} + 4 \right) \\ &= 4 \sum_{i=1}^k \Delta_i + 32 \sum_{i: \Delta_i > 0} \frac{\log(n)}{\Delta_i}. \end{aligned}$$

(b) Slightly larger than UCB, they are differ by scalar factor.

(c)

$$R(n) \lesssim 4 \sum_{i=1}^k \Delta_i + 16\alpha \sum_{i: \Delta_i > 0} \frac{\log(n)}{\Delta_i}.$$

(d) Refer "./Implementations"

Theorem (Cramér). Suppose $X_1, \dots, X_n$ are i.i.d. random variables whose values lie in $\mathbb{R}^d$ with finite empirical mean defined by $S_n = \frac{1}{n} \sum_{i=1}^n X_i$. Define μ_n to be the law of S_n . Then μ_n satisfies a large deviation principle on $\mathbb{R}^d$ with rate function given by the Legendre transform of the cumulant generating function of X_1 . That is,

- For any closed set $F \subset \mathbb{R}^d$,

$$\limsup_{n \rightarrow \infty} \frac{1}{n} \log \mu_n(F) \leq - \inf_{x \in F} \Lambda^*(x),$$

- For any open set $G \subset \mathbb{R}^d$,

$$\liminf_{n \rightarrow \infty} \frac{1}{n} \log \mu_n(G) \geq - \inf_{x \in G} \Lambda^*(x),$$

$$\text{where } \Lambda^*(x) = \sup_{\lambda \in \mathbb{R}^d} \{ \langle \lambda, x \rangle - \Lambda(\lambda) \} = \sup_{\lambda \in \mathbb{R}^d} \{ \langle \lambda, x \rangle - \log \mathbb{E} [e^{\langle \lambda, X_1 \rangle}] \}.$$

Theorem (Gärtner-Ellis). Suppose $X_1, \dots, X_n$ are random variables whose values lie in $\mathbb{R}^d$ with finite empirical mean defined by $S_n = \frac{1}{n} \sum_{i=1}^n X_i$. Define μ_n to be the law of S_n . Suppose the limit

$$\Lambda(\lambda) := \lim_{n \rightarrow \infty} \frac{1}{n} \log \mathbb{E} [e^{n \langle \lambda, S_n \rangle}]$$

exists for all $\lambda \in \mathbb{R}^d$ and $0 \in \mathcal{D}_\Lambda^\circ$, where $\mathcal{D}_\Lambda$ is the effective domain of Λ . Then

- For any closed set $F \subset \mathbb{R}^d$,

$$\limsup_{n \rightarrow \infty} \frac{1}{n} \log \mu_n(F) \leq - \inf_{x \in F} \Lambda^*(x),$$

- For any open set $G \subset \mathbb{R}^d$,

$$\liminf_{n \rightarrow \infty} \frac{1}{n} \log \mu_n(G) \geq - \inf_{x \in G \cap \mathcal{F}} \Lambda^*(x),$$

where $\Lambda^*(x)$ is the rate function given by the Legendre transform of $\Lambda(\lambda)$, and $\mathcal{F}$ is the set of exposed points of $\Lambda^*(x)$.

Furthermore, if Λ is essentially smooth and lower semicontinuous, then $\mathcal{F}$ can be replaced by the entire domain of Λ^* (i.e. LDP holds).